

Prometheus监控Kubernetes系列2——监控部署

原创： 罗佳豪 ServiceMesher 昨天



作者：罗佳豪 审校：孙海洲、李征、吴钧泽

这是该系列连载的第二篇，该系列文章：

[Prometheus监控Kubernetes系列1——监控框架](#)

由于容器化和微服务的大力发展，Kubernetes基本已经统一了容器管理方案，当我们使用Kubernetes来进行容器化管理的时候，全面监控Kubernetes也就成了我们第一个需要探索的问题。我们需要监控kubernetes的ingress、service、deployment、pod.....等等服务，以达到随时掌握Kubernetes集群的内部状况。此文章是Prometheus监控系列的第二篇，基于上一篇讲解了怎么对Kubernetes集群实施Prometheus监控。

K8s编排文件可参考 <https://github.com/xianyuLuo/prometheus-monitor-kubernetes>

Prometheus部署

在k8s上部署Prometheus十分简单，下面给的例子中将Prometheus部署到prometheus命名空间。

部署—数据采集

将kube-state-metrics和prometheus分开部署，先部署prometheus。

Prometheus

prometheus-rbac.yaml

```
1. apiVersion: rbac.authorization.k8s.io/v1beta1
2. kind: ClusterRole
3. metadata:
4.   name: prometheus
5. rules:
6. - apiGroups: [""]
7.   resources:
8.     - nodes
9.     - nodes/proxy
10.    - services
11.    - endpoints
12.    - pods
13.    verbs: ["get", "list", "watch"]
14. - apiGroups:
15.     - extensions
16.   resources:
17.     - ingresses
18.    verbs: ["get", "list", "watch"]
19. - nonResourceURLs: ["/metrics"]
20.   verbs: ["get"]
21. ---
22. apiVersion: v1
23. kind: ServiceAccount
24. metadata:
25.   name: prometheus
26.   namespace: prometheus
27. ---
28. apiVersion: rbac.authorization.k8s.io/v1beta1
29. kind: ClusterRoleBinding
30. metadata:
31.   name: prometheus
32. roleRef:
```

```
33.   apiGroup: rbac.authorization.k8s.io
34.   kind: ClusterRole
35.   name: prometheus
36. subjects:
37. - kind: ServiceAccount
38.   name: prometheus
39.   namespace: prometheus
```

prometheus.rbac.yml定义了Prometheus容器访问k8s apiserver所需的ServiceAccount、ClusterRole以及ClusterRoleBinding。

prometheus-config-configmap.yaml

```
1. apiVersion: v1
2. kind: ConfigMap
3. metadata:
4.   name: prometheus-config
5.   namespace: prometheus
6. data:
7.   prometheus.yml: |
8.     global:
9.       scrape_interval: 15s
10.      evaluation_interval: 15s
11.      scrape_configs:
12.
13.      - job_name: 'kubernetes-apisservers'
14.        kubernetes_sd_configs:
15.        - role: endpoints
16.        scheme: https
17.        tls_config:
18.          ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
19.          bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
20.        relabel_configs:
21.        - source_labels: [__meta_kubernetes_namespace,
22.          __meta_kubernetes_service_name, __meta_kubernetes_endpoint_port_name]
23.          action: keep
24.          regex: default;kubernetes;https
25.
26.      - job_name: 'kubernetes-nodes'
```

```
26.     kubernetes_sd_configs:
27.     - role: node
28.     scheme: https
29.     tls_config:
30.       ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
31.       bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
32.     relabel_configs:
33.     - action: labelmap
34.       regex: __meta_kubernetes_node_label_(.+)
35.     - target_label: __address__
36.       replacement: kubernetes.default.svc:443
37.     - source_labels: [__meta_kubernetes_node_name]
38.       regex: (.+)
39.       target_label: __metrics_path__
40.       replacement: /api/v1/nodes/${1}/proxy/metrics
41.
42.     - job_name: 'kubernetes-cadvisor'
43.     kubernetes_sd_configs:
44.     - role: node
45.     scheme: https
46.     tls_config:
47.       ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
48.       bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
49.     relabel_configs:
50.     - action: labelmap
51.       regex: __meta_kubernetes_node_label_(.+)
52.     - target_label: __address__
53.       replacement: kubernetes.default.svc:443
54.     - source_labels: [__meta_kubernetes_node_name]
55.       regex: (.+)
56.       target_label: __metrics_path__
57.       replacement: /api/v1/nodes/${1}/proxy/metrics/cadvisor
58.
59.     - job_name: 'kubernetes-service-endpoints'
60.     kubernetes_sd_configs:
61.     - role: endpoints
62.     relabel_configs:
```

```
63.     - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_scrape]
64.         action: keep
65.         regex: true
66.     - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_scheme]
67.         action: replace
68.         target_label: __scheme__
69.         regex: (https?)
70.     - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_path]
71.         action: replace
72.         target_label: __metrics_path__
73.         regex: (.+)
74.     - source_labels: [__address__,
__meta_kubernetes_service_annotation_prometheus_io_port]
75.         action: replace
76.         target_label: __address__
77.         regex: ([^:]+)(?::\d+)?;(\d+)
78.         replacement: $1:$2
79.     - action: labelmap
80.         regex: __meta_kubernetes_service_label_(.+)
81.     - source_labels: [__meta_kubernetes_namespace]
82.         action: replace
83.         target_label: kubernetes_namespace
84.     - source_labels: [__meta_kubernetes_service_name]
85.         action: replace
86.         target_label: kubernetes_name
87.
88.     - job_name: 'kubernetes-services'
89.     kubernetes_sd_configs:
90.     - role: service
91.       metrics_path: /probe
92.       params:
93.         module: [http_2xx]
94.       relabel_configs:
```

```
95.     - source_labels:
[__meta_kubernetes_service_annotation_prometheus_io_probe]
96.         action: keep
97.         regex: true
98.     - source_labels: [__address__]
99.         target_label: __param_target
100.    - target_label: __address__
101.        replacement: blackbox-exporter.example.com:9115
102.    - source_labels: [__param_target]
103.        target_label: instance
104.    - action: labelmap
105.        regex: __meta_kubernetes_service_label_(.+)
106.    - source_labels: [__meta_kubernetes_namespace]
107.        target_label: kubernetes_namespace
108.    - source_labels: [__meta_kubernetes_service_name]
109.        target_label: kubernetes_name
110.
111.    - job_name: 'kubernetes-ingresses'
112.      kubernetes_sd_configs:
113.      - role: ingress
114.      relabel_configs:
115.      - source_labels:
[__meta_kubernetes_ingress_annotation_prometheus_io_probe]
116.          action: keep
117.          regex: true
118.      - source_labels:
[__meta_kubernetes_ingress_scheme,__address__,__meta_kubernetes_ingress_path]
119.          regex: (.+);(.+);(.+)
120.          replacement: ${1}://${2}${3}
121.          target_label: __param_target
122.      - target_label: __address__
123.          replacement: blackbox-exporter.example.com:9115
124.      - source_labels: [__param_target]
125.          target_label: instance
126.      - action: labelmap
127.          regex: __meta_kubernetes_ingress_label_(.+)
128.      - source_labels: [__meta_kubernetes_namespace]
```

```

129.     target_label: kubernetes_namespace
130.   - source_labels: [__meta_kubernetes_ingress_name]
131.     target_label: kubernetes_name
132.
133.   - job_name: 'kubernetes-pods'
134.     kubernetes_sd_configs:
135.       - role: pod
136.     relabel_configs:
137.       - source_labels:
138.         [__meta_kubernetes_pod_annotation_prometheus_io_scrape]
139.         action: keep
140.         regex: true
141.       - source_labels:
142.         [__meta_kubernetes_pod_annotation_prometheus_io_path]
143.         action: replace
144.         target_label: __metrics_path__
145.         regex: (.+)
146.       - source_labels: [__address__,
147.         __meta_kubernetes_pod_annotation_prometheus_io_port]
148.         action: replace
149.         regex: ([^:]+)(?::\d+)?;(\d+)
150.         replacement: $1:$2
151.         target_label: __address__
152.       - action: labelmap
153.         regex: __meta_kubernetes_pod_label_(.+)
154.       - source_labels: [__meta_kubernetes_namespace]
155.         action: replace
156.         target_label: kubernetes_namespace

```

prometheus-config-configmap.yaml定义了prometheus的配置文件，以configmap的形式使用。

prometheus-dep.yaml

1. apiVersion: apps/v1beta2
2. kind: Deployment

```
3. metadata:
4.   name: prometheus-dep
5.   namespace: prometheus
6. spec:
7.   replicas: 1
8.   selector:
9.     matchLabels:
10.      app: prometheus-dep
11.   template:
12.     metadata:
13.       labels:
14.         app: prometheus-dep
15.       spec:
16.         containers:
17.           - image: prom/prometheus:v2.3.2
18.             name: prometheus
19.             command:
20.               - "/bin/prometheus"
21.             args:
22.               - "--config.file=/etc/prometheus/prometheus.yml"
23.               - "--storage.tsdb.path=/prometheus"
24.               - "--storage.tsdb.retention=1d"
25.             ports:
26.               - containerPort: 9090
27.                 protocol: TCP
28.             volumeMounts:
29.               - mountPath: "/prometheus"
30.                 name: data
31.               - mountPath: "/etc/prometheus"
32.                 name: config-volume
33.             resources:
34.               requests:
35.                 cpu: 100m
36.                 memory: 100Mi
37.               limits:
38.                 cpu: 500m
39.                 memory: 2500Mi
40.             serviceAccountName: prometheus
```



```
41.   imagePullSecrets:
42.     - name: regsecret
43.   volumes:
44.     - name: data
45.       emptyDir: {}
46.     - name: config-volume
47.       configMap:
48.         name: prometheus-config
```

prometheus-dep.yaml定义了prometheus的部署，这里使用--storage.tsdb.retention参数，监控数据只保留1天，因为最终监控数据会统一汇总。

limits资源限制根据集群大小进行适当调整。

prometheus-svc.yaml

```
1. kind: Service
2. apiVersion: v1
3. metadata:
4.   name: prometheus-svc
5.   namespace: prometheus
6. spec:
7.   type: NodePort
8.   ports:
9.     - port: 9090
10.       targetPort: 9090
11.       nodePort: 30090
12.   selector:
13.     app: prometheus-dep
```

prometheus-svc.yaml定义Prometheus的Service，需要将Prometheus以NodePort、LoadBalancer或Ingress暴露到集群外部，这样外部的Prometheus才能访问它。这里采用的NodePort，所以只需要访问集群中有外网地址的任意一台服务器的30090端口就可以使用prometheus。

kube-state-metrics

prometheus部署成功后，接着再部署kube-state-metrics作为prometheus的一个exporter来使用，提供deployment、daemonset、cronjob等服务的监控数据。

kube-state-metrics-rbac.yaml

```
1. apiVersion: v1
2. kind: ServiceAccount
3. metadata:
4.   name: kube-state-metrics
5.   namespace: prometheus
6. ---
7.
8. apiVersion: rbac.authorization.k8s.io/v1
9. # kubernetes versions before 1.8.0 should use
rbac.authorization.k8s.io/v1beta1
10. kind: Role
11. metadata:
12.   namespace: prometheus
13.   name: kube-state-metrics-resizer
14. rules:
15. - apiGroups: [""]
16.   resources:
17.   - pods
18.   verbs: ["get"]
19. - apiGroups: ["extensions"]
20.   resources:
21.   - deployments
22.   resourceNames: ["kube-state-metrics"]
23.   verbs: ["get", "update"]
24. ---
25.
26. apiVersion: rbac.authorization.k8s.io/v1
27. # kubernetes versions before 1.8.0 should use
rbac.authorization.k8s.io/v1beta1
28. kind: ClusterRole
29. metadata:
30.   name: kube-state-metrics
31. rules:
32. - apiGroups: [""]
33.   resources:
34.   - configmaps
35.   - secrets
```

```
36.   - nodes
37.   - pods
38.   - services
39.   - resourcequotas
40.   - replicationcontrollers
41.   - limitranges
42.   - persistentvolumeclaims
43.   - persistentvolumes
44.   - namespaces
45.   - endpoints
46.   verbs: ["list", "watch"]
47. - apiGroups: ["extensions"]
48.   resources:
49.   - daemonsets
50.   - deployments
51.   - replicaset
52.   verbs: ["list", "watch"]
53. - apiGroups: ["apps"]
54.   resources:
55.   - statefulsets
56.   verbs: ["list", "watch"]
57. - apiGroups: ["batch"]
58.   resources:
59.   - cronjobs
60.   - jobs
61.   verbs: ["list", "watch"]
62. - apiGroups: ["autoscaling"]
63.   resources:
64.   - horizontalpodautoscalers
65.   verbs: ["list", "watch"]
66. ---
67.
68. apiVersion: rbac.authorization.k8s.io/v1
69. # kubernetes versions before 1.8.0 should use
rbac.authorization.k8s.io/v1beta1
70. kind: RoleBinding
71. metadata:
```

```

72.   name: kube-state-metrics
73.   namespace: prometheus
74. roleRef:
75.   apiGroup: rbac.authorization.k8s.io
76.   kind: Role
77.   name: kube-state-metrics-resizer
78. subjects:
79. - kind: ServiceAccount
80.   name: kube-state-metrics
81.   namespace: prometheus
82. ---
83.
84. apiVersion: rbac.authorization.k8s.io/v1
85. # kubernetes versions before 1.8.0 should use
rbac.authorization.k8s.io/v1beta1
86. kind: ClusterRoleBinding
87. metadata:
88.   name: kube-state-metrics
89. roleRef:
90.   apiGroup: rbac.authorization.k8s.io
91.   kind: ClusterRole
92.   name: kube-state-metrics
93. subjects:
94. - kind: ServiceAccount
95.   name: kube-state-metrics
96.   namespace: prometheus

```

kube-state-metrics-rbac.yaml定义了kube-state-metrics访问k8s apiserver所需的ServiceAccount和ClusterRole及ClusterRoleBinding。

kube-state-metrics-dep.yaml

```

1. apiVersion: apps/v1beta2
2. # Kubernetes versions after 1.9.0 should use apps/v1
3. # Kubernetes versions before 1.8.0 should use apps/v1beta1 or
extensions/v1beta1
4. # addon-resizer描述:
https://github.com/kubernetes/autoscaler/tree/master/addon-resizer
5. kind: Deployment

```

```
6. metadata:
7.   name: kube-state-metrics
8.   namespace: prometheus
9. spec:
10.  selector:
11.    matchLabels:
12.      k8s-app: kube-state-metrics
13.  replicas: 1
14.  template:
15.    metadata:
16.      labels:
17.        k8s-app: kube-state-metrics
18.    spec:
19.      serviceAccountName: kube-state-metrics
20.      containers:
21.        - name: kube-state-metrics
22.          image: xianyuluo/kube-state-metrics:v1.3.1
23.          ports:
24.            - name: http-metrics
25.              containerPort: 8080
26.            - name: telemetry
27.              containerPort: 8081
28.          readinessProbe:
29.            httpGet:
30.              path: /healthz
31.              port: 8080
32.              initialDelaySeconds: 5
33.              timeoutSeconds: 5
34.            - name: addon-resizer
35.              image: xianyuluo/addon-resizer:1.7
36.          resources:
37.            limits:
38.              cpu: 100m
39.              memory: 30Mi
40.            requests:
41.              cpu: 100m
42.              memory: 30Mi
43.          env:
```

```
44.     - name: MY_POD_NAME
45.       valueFrom:
46.         fieldRef:
47.           fieldPath: metadata.name
48.     - name: MY_POD_NAMESPACE
49.       valueFrom:
50.         fieldRef:
51.           fieldPath: metadata.namespace
52.     command:
53.     - /pod_nanny
54.     - --container=kube-state-metrics
55.     - --cpu=100m
56.     - --extra-cpu=1m
57.     - --memory=100Mi
58.     - --extra-memory=2Mi
59.     - --threshold=5
60.     - --deployment=kube-state-metrics
```

kube-state-metrics-dep.yaml定义了kube-state-metrics的部署。

kube-state-metrics-svc.yaml

```
1. apiVersion: v1
2. kind: Service
3. metadata:
4.   name: kube-state-metrics
5.   namespace: prometheus
6.   labels:
7.     k8s-app: kube-state-metrics
8.   annotations:
9.     prometheus.io/scrape: 'true'
10. spec:
11.   ports:
12.     - name: http-metrics
13.       port: 8080
14.       targetPort: http-metrics
15.       protocol: TCP
16.     - name: telemetry
17.       port: 8081
18.       targetPort: telemetry
```

19. `protocol: TCP`
20. `selector:`
21. `k8s-app: kube-state-metrics`

kube-state-metrics-svc.yaml定义了kube-state-metrics的暴露方式，这里只需要使用默认的ClusterIP就可以了，因为它只提供给集群内部的promethes访问。

k8s集群中的prometheus监控到这儿就已经全部OK了，接下来还需要做的是汇总数据、展示数据及告警规则配置。

部署—数据汇总

prometheus-server

prometheus-server和前面prometheus的步骤基本相同，需要针对configmap、数据存储时间（一般为30d）、svc类型做些许改变，同时增加 rule.yaml。

prometheus-server不需要kube-state-metrics。prometheus-server可以部署在任意k8s集群，或者部署在K8s集群外部都可以。

prometheus-rbac.yaml (内容和上面的一致，namespace为prometheus-server)

1.

prometheus-server-config-configmap.yaml

1. `apiVersion: v1`
2. `kind: ConfigMap`
3. `metadata:`
4. `name: prometheus-server-config`
5. `namespace: prometheus-server`
6. `data:`
7. `prometheus.yml: |`
8. `global:`
9. `scrape_interval: 30s # Set the scrape interval to every 15 seconds.`
`Default is every 1 minute.`
10. `evaluation_interval: 30s # Evaluate rules every 15 seconds. The`
`default is every 1 minute.`
11. `scrape_timeout: 30s`
- 12.
13. `# Alertmanager configuration`
14. `alerting:`
15. `alertmanagers:`
16. `- static_configs:`

```

17.     - targets: ["x.x.com:80"]
18.
19.   rule_files:
20.     - "/etc/prometheus/rule.yml"
21.
22.   scrape_configs:
23.     - job_name: 'federate-k8scluster-1'
24.       scrape_interval: 30s
25.       honor_labels: true
26.       metrics_path: '/federate'
27.       params:
28.         'match[]':
29.           - '{job=~"kubernetes-.*"}'
30.       static_configs:
31.         - targets: ['x.x.x.x:30090']
32.         labels:
33.           k8scluster: xxxx-k8s
34.     - job_name: 'federate-k8scluster-2'
35.       scrape_interval: 30s
36.       honor_labels: true
37.       metrics_path: '/federate'
38.       params:
39.         'match[]':
40.           - '{job=~"kubernetes-.*"}'
41.       static_configs:
42.         - targets: ['x.x.x.x:30090']
43.         labels:
44.           k8scluster: yyyy-k8s

```

global：全局配置。设置了收集数据频率、超时等

alerting：告警配置。指定了prometheus将满足告警规则的信息发送到哪儿？告警规则在rule_files定义

rule_files：定义的告警规则文件

scrape_configs：监控数据刮取配置。定义了2个job，分别是federate-k8scluster-1、federate-k8scluster-2。其中federate-k8scluster-1配置了去x.x.x.x30090采集数据，并且要匹配job名为"kubernetes-"开头。注意下面的

labels, 这个是自己定义的, 它的作用在于给每一条刮取过来的监控数据都加上一个 **k8scluster: xxxx-k8s** 的Key-Value, xxxx一般指定为项目代码。这样我们可以在多个集群数据中区分该条数据是属于哪一个k8s集群, 这对于后面的展示和告警都非常有利。

prometheus-server-rule-configmap.yaml

```
1. apiVersion: v1
2. kind: ConfigMap
3. metadata:
4.   name: prometheus-server-rule-config
5.   namespace: prometheus-server
6. data:
7.   rule.yml: |
8.     groups:
9.     - name: kubernetes
10.       rules:
11.       - alert: PodDown
12.         expr: kube_pod_status_phase{phase="Unknown"} == 1 or
kube_pod_status_phase{phase="Failed"} == 1
13.         for: 1m
14.         labels:
15.           severity: error
16.           service: prometheus_bot
17.           receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
18.         annotations:
19.           summary: Pod Down
20.           k8scluster: "{{ $labels.k8scluster }}"
21.           namespace: "{{ $labels.namespace }}"
22.           pod: "{{ $labels.pod }}"
23.           container: "{{ $labels.container }}"
24.
25.       - alert: PodRestart
26.         expr: changes(kube_pod_container_status_restarts_total{pod !~
"analyzer.*"}[10m]) > 0
27.         for: 1m
28.         labels:
29.           severity: error
```

```

30.     service: prometheus_bot
31.     receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
32.     annotations:
33.         summary: Pod Restart
34.         k8scluster: "{{ $labels.k8scluster }}"
35.         namespace: "{{ $labels.namespace }}"
36.         pod: "{{ $labels.pod }}"
37.         container: "{{ $labels.container }}"
38.
39.     - alert: NodeUnschedulable
40.         expr: kube_node_spec_unschedulable == 1
41.         for: 5m
42.         labels:
43.             severity: error
44.             service: prometheus_bot
45.             receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
46.             annotations:
47.                 summary: Node Unschedulable
48.                 k8scluster: "{{ $labels.k8scluster }}"
49.                 node: "{{ $labels.node }}"
50.
51.     - alert: NodeStatusError
52.         expr: kube_node_status_condition{condition="Ready", status!="true"}
53.         == 1
54.         for: 5m
55.         labels:
56.             severity: error
57.             service: prometheus_bot
58.             receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
59.             annotations:
60.                 summary: Node Status Error
61.                 k8scluster: "{{ $labels.k8scluster }}"
62.                 node: "{{ $labels.node }}"
63.
64.     - alert: DaemonsetUnavailable
65.         expr: kube_daemonset_status_number_unavailable > 0
66.         for: 5m

```

```

66.     labels:
67.         severity: error
68.         service: prometheus_bot
69.         receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
70.     annotations:
71.         summary: Daemonset Unavailable
72.         k8scluster: "{{ $labels.k8scluster }}"
73.         namespace: "{{ $labels.namespace }}"
74.         daemonset: "{{ $labels.daemonset }}"
75.
76.     - alert: JobFailed
77.       expr: kube_job_status_failed == 1
78.       for: 5m
79.       labels:
80.         severity: error
81.         service: prometheus_bot
82.         receiver_group: "{{ $labels.k8scluster }}_{{ $labels.namespace }}"
83.     annotations:
84.         summary: Job Failed
85.         k8scluster: "{{ $labels.k8scluster }}"
86.         namespace: "{{ $labels.namespace }}"
87.         job: "{{ $labels.exported_job }}"

```

rule.yaml定义了告警规则。此文件中定义了 PodDown、PodRestart、NodeUnschedulable、NodeStatusError、DaemonsetUnavailable、JobFailed 共6条规则。

alert: 名称

expr: 表达式。prometheus的SQL语句

for: 时间范围

annotations: 告警消息，其中 {{*}} 为Prometheus内部变量

prometheus-server-dep.yaml (参考上面的prometheus-dep.yaml做些许调整)

```

1. apiVersion: apps/v1beta2
2. kind: Deployment
3. metadata:
4.   name: prometheus-server-dep
5.   namespace: prometheus-server

```

```

6. spec:
7.   replicas: 1
8.   selector:
9.     matchLabels:
10.      app: prometheus-server-dep
11.      .....
12.     args:
13.      - "--config.file=/etc/prometheus/prometheus.yml"
14.      - "--storage.tsdb.path=/prometheus"
15.      - "--web.console.libraries=/usr/share/prometheus/console_libraries"
16.      - "--web.console.templates=/usr/share/prometheus"
17.      - "--storage.tsdb.retention=30d"
18.      - "--web.enable-lifecycle"
19.      .....
20.     volumeMounts:
21.      - mountPath: "/prometheus"
22.        name: data
23.      - mountPath: "/etc/prometheus/prometheus.yml"
24.        name: server-config-volume
25.        subPath: prometheus.yml
26.      - mountPath: "/etc/prometheus/rule.yml"
27.        name: rule-config-volume
28.        subPath: rule.yml
29.      .....
30.     volumes:
31.      - name: data
32.        emptyDir: {}
33.      - name: server-config-volume
34.        configMap:
35.          name: prometheus-server-config
36.      - name: rule-config-volume
37.        configMap:
38.          name: prometheus-server-rule-config

```

volumes.data这里使用的是emptyDir，这样其实不妥，应该单独挂载一块盘来存储汇总数据。可使用pv实现。

prometheus-server-svc.yaml (参考上面的prometheus-svc.yaml做些许调整)

```

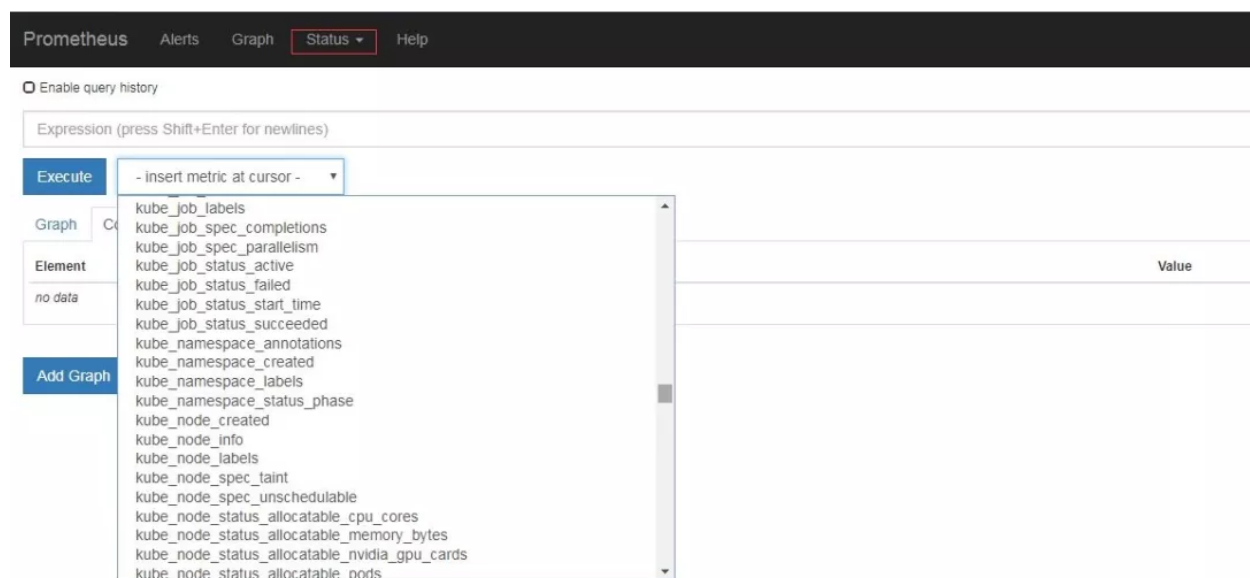
1. kind: Service

```

2. `apiVersion: v1`
3. `metadata:`
4. `name: prometheus-server-svc`
5. `namespace: prometheus-server`
6. `spec:`
7. `type: LoadBalancer`
8. `ports:`
9. - `port: 80`
10. `targetPort: 9090`
11. `selector:`
12. `app: prometheus-server-dep`

到这儿，数据采集和数据汇总就已经OK了。

Prometheus-server部署成功之后，在浏览器中可以看到监控数据汇总信息了



Status --> Configuration 中可以看到Prometheus-server的配置

Status --> Rules 中可以看到规则文件内容

Status --> Targets 中可以看到刮取目标的状态信息

告警配置

遵循上篇文章中的架构，告警使用Prometheus官方提供的组件Alertmanager

alertmanager-config-configmap.yaml

```
1. apiVersion: v1
2. kind: ConfigMap
3. metadata:
4.   name: alertmanager-config
5.   namespace: prometheus
6. data:
7.   config.yml: |
8.     global:
9.       resolve_timeout: 5m
10.
11.     route:
12.       receiver: default
13.       group_wait: 30s
14.       group_interval: 5m
15.       repeat_interval: 4h
16.       group_by: ['alertname', 'k8scluster', 'node', 'container',
'exported_job', 'daemonset']
17.     routes:
18.     - receiver: send_msg_warning
19.       group_wait: 60s
20.       match:
21.         severity: warning
22.
23.     receivers:
24.     - name: default
25.       webhook_configs:
26.       - url: 'http://msg.x.com/xxx/'
27.         send_resolved: true
28.       http_config:
29.         bearer_token: 'xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx'
30.
31.     - name: send_msg_warning
32.       webhook_configs:
33.       - url: 'http://msg.x.com/xxx/'
34.         send_resolved: true
```

```
35. http_config:
36. bearer_token: 'xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx'
```

alertmanager-config-configmap.yaml定义了alertmanager的配置文件

route: 路由。分级匹配, 然后交给指定 receivers, 其中route.group_by中的

k8scluster是prometheus-server-config.yaml中自定义的标签

receivers: 发送。这里使用webhook方式发送给自研的send_msg模块

email、wechat、webhook、slack等发送方式配置请见官网文档:

<https://prometheus.io/docs/alerting/configuration/>

alertmanager-dep.yaml

```

1. apiVersion: apps/v1beta2
2. kind: Deployment
3. metadata:
4.   name: alertmanager-dep
5.   namespace: alertmanager
6. spec:
7.   replicas: 1
8.   selector:
9.     matchLabels:
10.      app: alertmanager-dep
11.   template:
12.     metadata:
13.       labels:
14.         app: alertmanager-dep
15.     spec:
16.       containers:
17.       - image: prom/alertmanager:v0.15.2
18.         name: alertmanager
19.         args:
20.         - "--config.file=/etc/alertmanager/config.yml"
21.         - "--storage.path=/alertmanager"
22.         - "--data.retention=720h"
23.       volumeMounts:
24.       - mountPath: "/alertmanager"
25.         name: data
26.       - mountPath: "/etc/alertmanager"

```

```

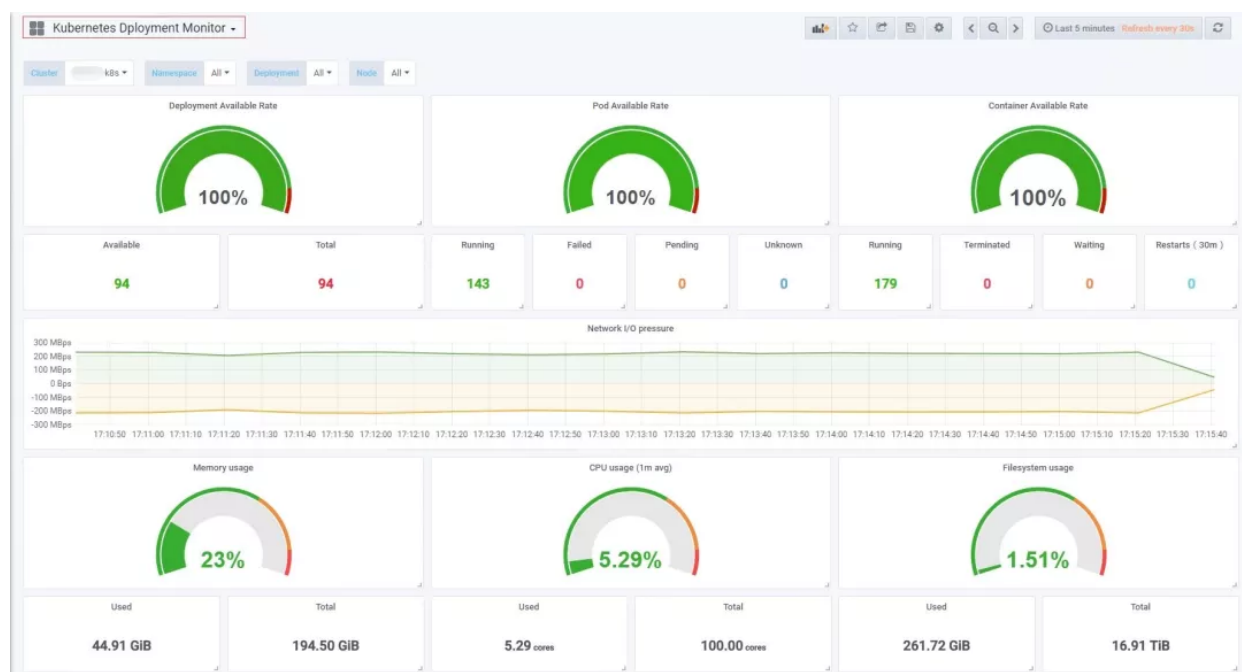
27.     name: config-volume
28.   resources:
29.     requests:
30.       cpu: 100m
31.       memory: 100Mi
32.     limits:
33.       cpu: 500m
34.       memory: 2500Mi
35.   volumes:
36.   - name: data
37.     emptyDir: {}
38.   - name: config-volume
39.     configMap:
40.       name: alertmanager-config

```

alertmanager-dep.yaml定义了Alertmanager的部署。

展示

遵循上篇文章中的架构，展示使用开源的Grafana。Grafana的部署方式就不详细描述了，下面展示两个Dashboard



kubernetes-deployment-dashboard，展示了大多关于deployment的信息。左上角的Cluster选项就是利用prometheus-server-config.yaml中自定义的labels.k8scluster标签实现的。



kubernetes-pod-dashboard，展示的都是关于pod和container的信息，包括CPU、mem使用监控。此页面数据量一般比较大。左上角的Cluster选项也是利用prometheus-server-config.yaml中自定义的labels.k8scluster做的。

kubernetes-deployment-dashboard下载地址：

<https://grafana.com/dashboards/9730>

kubernetes-pod-dashboard下载地址：<https://grafana.com/dashboards/9729>

结束

详细监控Kubernetes集群本身就是一项复杂的工作，好在有Prometheus、Grafana、kube-state-metrics这些优秀的开源工具，才让我们的工作复杂度得以缓解，Thanks。

此文章也是“使用prometheus完美监控kubernetes集群”系列的第二篇，如果在部署过程中遇到问题或者有不理解的地方，欢迎随时后台留言。

推荐阅读

[Istio中使用Prometheus进行应用程序指标度量](#)

[全手动部署prometheus operator监控kubernetes集群以及一些坑](#)

[Envoy service mesh、Prometheus和Grafana下的微服务监控](#)

[对使用Kubernetes和Istio管理的基于容器的基础设施进行全面服务监控](#)